



ross 2

Date	@October 8, 2025
Category	
Summary	The team has developed a query system with Pinecone integration and a database for client use, including feedback capabilities. The system categorizes queries into three main categories, each with specific questions, but faces challenges in correctly identifying categories and managing project timelines. Technical components include using string enums for deterministic answers and a need for Lang Graph implementation. Action items include investigating the Lang Graph framework and implementing the question pipeline.

Summary

Current System Status

- Team has implemented a query system with Pinecone integration, initially containing 200-250 parts
- A database with files, docs, and PDFs has been created for client use
- System includes feedback capability where AI responses can be edited and saved in both Pinecone and MongoDB
- Demo has previously been given to Jeff with live feedback capability

Taxonomy and Pipeline Architecture

- System needs to categorize queries before asking specific questions
- Current plan involves three main categories (including "clutch" and "brakes")
- Each category will have its own set of questions

- Team is developing a pipeline where:
 - First step identifies the correct category
 - Second step asks category-specific questions
 - System validates user answers against expected data types
 - Final output adds all query attributes

Implementation Challenges

- Concern about correctly identifying the initial category, as choosing wrong category would invoke incorrect question pipeline
- Discussion about vertical scaling challenges if subcategories are added beneath main categories
- Timeline pressure with 6-month project timeline (1 month already passed) and January 1st deadline
- Debate about whether to wait for complete taxonomy before proceeding

Technical Components

- System uses string enums for deterministic answers rather than free text
- User query attributes will be used to create appropriate responses
- Comparison made to car recommendation system (e.g., recommending XUV700 to replace XUV500)
- Lang Graph implementation needed for the pipeline architecture

Action Items

- ☐ Shashank to investigate Lang Graph framework (estimated to take 1-2 days)
- ☐ Shashank to review "small projects" and JSON output of the identical parser
- ☐ Team to implement question pipeline with the three existing categories
- ☐ Someone to connect with Shashank after the call to provide detailed instructions

☐ Consider putting a request in the ROS group tagging Kaushik regarding "Tech Tsunami" if needed

Notes

Agenda

What are the topics to be discussed?

Notes

Add any notes to provide additional context and details

Transcript

Obviously, there used to be a lot of parts in it. I mean, when I did the query, even after the query, there were 200-250 parts. I have set the question in such a way that the answer is correct. So, that's how I did it. I did everything on Pinecone. Okay. So, what you're doing, the metadata search, metadata SMS search, those were the chunks.

I don't know how to do it. I don't know how to do it.

files and docs and pdfs and we made a mock database for the database This is the actual database Yes, we made our own database for them Is there any content or not? Sir, I mean, it was a long time ago, we will have to see. Right now, it is not there for some reason. Right now, it is files and docs and PDFs.

So what is the difference between... See, he gave this live. At the time of the demo. This is live feedback from Jeff. You can also edit the response given by the AI. Click on edit. It's a little buggy. If you click on it, it will look like GitHub. Please remove the line below for long service life and reliability. Please remove it. So this will be done like this.

Yeah, I saved it. Next time when he asks for the query, this should be the response, right? Yes, it is being saved in the system. So, when we get the response, it will be understood. This exact response will not come. But if you have any feedback, it is saved in both Pinecone and Mongo.

What happens in this? Open analysis. Completed user analysis. In this, we catch everything.

I'm going to be a provide more I'm going to be a provide more I'm going to be a provide more information.

Yes, so sir, every category will have a question set, right? Every category will have different questions. Even if you search, you will find different questions. So, the first thing is, the one for whom we are making this, is a junior sales engineer, so he won't be able to find the category by himself.

So, we have to find the category first before asking questions, right? So, we have to find the category first before making a note of the line graph. It will be like a key. You can save it in Mongo or use a hardcode for now. After that, it will go to questions.

It will be either enum or query, so my query will be on a deterministic set. The user can't write anything in this answer.

For example, if I am asking a question about the temperature of the system, I have written the answer in it. You have to do it in the string enum. Chang bhai, did you understand? The string enum is what is written there. It has been edited to include proper punctuation.

It's already trained. You'll check it, right? A lot of things in ROS are already trained. Sir, can you open the tab of ROS website?

Hmm. Yeah, that was for class B control.

It has been edited to include proper punctuation.

That's weird.

I'm not sure if you can hear me.

Thank you very much.

Okay. Okay. In this, this one, click on this one, it has 261 parts. Okay.

Here we have a missed request of Ketanath Tanesh, want to update this crash break circuit and automatic counter balance and replace our main air lockout system, what do we do?

In this, the final output will be that all the attributes of the query have been added. So, in Mongo, it will be... Shashank bhai? Hello? Yes, tell me.

It will be incomplete. Once the whole taxonomy is discussed and the back end is clear, only then can we go there. Because the first thing the question has to look at is the category, right? I think it's going to be I think it's going to be I think it's going to be right.

And you said in the note that there will be a category out of these 10 categories, right? That can be a horizontal scale. There is no problem with that. I just have to add it in the prompt. But if there is a problem with the vertical scaling, suppose there is a game even after a category or subcategory, then...

The first thing I want to say is that if I define a category wrongly or if I read the question wrongly and the question was about the clutch, I have asked a few more questions and due to that, a new question pipeline will be invoked. So, that is a problem. Yes, that is a problem.

For example, if we talk about a normal day, I went to a showroom between a SUV and a sedan. I mistook the customer for being interested in the SUV and started asking him about the sedan. So, our first category to pick is very important. In that, I am saying that if we are picking the right category,

This is a one-layer category, and you are saying that there is nothing below the three clutch brakes, so we are good.

But we can't wait, you know, But we can't wait, you know, But we can't wait, you know, a state will be confirmed. Sir, it should be the same. I don't know, sir. It will be done. Sir, let me tell you something very clear. The same thing happened to me in Star Mortgage. Same to same. It's Shashank bhai. It's a matter of identity. It's not a big deal. It was the same recommendation. It happened in the same story.

It has been edited to include proper punctuation. It has been edited to include proper punctuation.

We will have to go to the taxonomy in the next call. Please close the taxonomy and close the questions. Otherwise, we will make it and your efforts will be added on. But we cannot say that because this is a 6-month project and a month has already passed.

Now, because we have to give it on 1st January. Do you understand what I am saying? If these things are added vertically, then sir, the effort increases a lot. I am not just saying that. I appreciate all your pushback. But, for now, the current story, go with the current state. Right now, we are not in a position to push back. The direction is...

We have these three categories. Take a try to build the. Problem is this pipeline. Again, setting Hawaii will be basically. Yeah, yeah. Again, I'll be. We do not have that, uh, I would say a leverage, uh, space. We wait for a week or we'll go back to Kaushik and say once we connect with them we'll get that clarity with taxonomy and all those things. We don't have that space.

It is not it is not. we can make it but like this we can make it but like this we can make it but like this is this is not like. and then we should have created and then we should have created This is how it should be. He taught us this. In Agent X, we were making something out of rag. And we were doing something out of rag. Then he said we should find entities and call UDL.

And from there we understood that recommendation always comes first by getting the taxonomy right, then to the questions.

How many levels of taxonomy are there? That's the thing, sir. But, okay. If you say so, we can make it. But it would be a rework. There is an erection on top. I'm going to have to help you. I'm going to have to help you. I'm going to have to help you. because uh

It has been edited to include proper punctuation.

No, no, bro. Don't worry. It's not that tough. There is a resource here. You can utilize it. It's possible if you want to give it by next week. We can all use it. You can do that. It is a little laggy. And it is not that easy. Yes, it was in it. Sir, should I put it in the group, in the Ross group, that by tagging Kausik, we need Tech Tsunami, if there is any issue. Sir, you are thinking that I am saying absolutely wrong, or I am overexaggerating this thing too much, this can also happen.

I'll do it. I'll connect with Shashank after the call. Shashank, I'll call you back. I'll explain it to you later. What you have to do is They do the work, you can just monitor So that they learn at the same time How to do it, what to do With you, Dhanpal and Sashank You have to keep up with the monitoring activity

I forgot what I was going to say. You can think of it as buying a car, right? So, if a user comes to you and posts a message in the group saying I had an XUV500 car. Now it is not working fine. Please suggest me a similar car like it. So, Mahindra shut it down. Now, there is a similar car, the XUV700. So, Mahindra will recommend the XUV700. So, basically, what is the scene in ROS? Their team, their valves shut down.

Obviously, the life of a car is also important. When a car breaks down, we have to replace it with a new one. Understood? This is a sedan, this is an SUV, or whatever it is, right? Because you have to recommend it to the person in front of you. Okay? There will be different questions for Sedans and SUVs. You know that you bought a car, right? The Sedan guy will ask different questions.

XUV Category The question is, does size matter in XUV? Does IR size matter in XUV? It will show you the part number that needs to be replaced. After that, it will show you the knowledge add-on. What is the size? What is the tire size? So, XUV500 and XUV Mahindra have Bolero as well.

It has been edited to include proper punctuation. It has been edited to include proper punctuation. So when you find a category, like Clutch, Category A, Category B, Category C, suppose you find B category, then the questions of B category will be very different from A category.

Now, the user has already answered some of the questions in the prompt. Like, I want this much temperature or this much pressure. So, you won't ask the same temperature and pressure questions again.

You will ask that question, and after that, when you ask that question, you will validate the answer that he will give, whether he has given the correct answer or whatever he has said.

It is not in the data. The answer to the question is that I asked the temperature in it. And it cannot say that it is a kilometer.

No, even that doesn't matter. If you show a null, then there is no recommendation. Change your preferences. That doesn't matter to you. You have to at least check what the user is saying. So, this is how you make a Shazang pipeline. This is the question and its answer and according to it, the output.

The game runs on the category level and the game runs on the attribute level. The attribute level is basically used to create queries.

Shashank, I need your help. Should I drop it? Yes, Sir. For now, Shashank will have to give some time for the lang graph. It will take at least a day or two to understand it, sir.

Shashank bhai, you can look at two things. You can look at Langraab once. And you can look at small projects. Don't go too deep. Look at it from the top. Take a basic idea of what it is. You can also see the JSON output of the identical parser.

I'll drop it. Okay, okay. Shashank bhai, you start from the line graph. I have to make a personal call.

It has been edited to include proper punctuation. I will call you and let you know. Okay, call me.

I